

# Model B: Imbalanced Review Rating Predictor

## Purpose

Model A is a deep learning-based text classification model built to predict review ratings (1 to 5 stars) from raw product review text. It was trained on a **balanced dataset**, where each rating class (1 through 5) has an equal number of samples. This design helps ensure fairness by preventing the model from being biased toward dominant classes.

## Full code:

```
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Embedding, LSTM, Dense, Bidirectional, Dropout
from tensorflow.keras.callbacks import EarlyStopping

model = Sequential()

model.add(Embedding(input_dim=vocab_size,
                    output_dim=embedding_dim,
                    weights=[embedding_matrix],
                    input_length=max_len,
                    trainable=True))
model.add(Bidirectional(LSTM(128, return_sequences=False)))
model.add(Dropout(0.5)) # optional but helps reduce overfitting
model.add(Dense(32, activation='relu'))
model.add(Dense(5, activation='softmax')) # 5 classes for ratings

model.compile(loss='sparse_categorical_crossentropy', optimizer='adam', metrics=['accuracy'])
model.summary()

# Train
early_stop = EarlyStopping(monitor='val_loss', patience=2, restore_best_weights=True)
history = model.fit(X_train_pad, y_train, epochs=20, batch_size=64, validation_data=(X_test_pad,
y_test), callbacks=[early_stop])
```

## Model Architecture

The model is constructed using the **Keras Sequential API** and follows a typical embedding + RNN + dense classifier pipeline. Here's a breakdown of each component:

### 1. Embedding Layer

```
Embedding(input_dim=vocab_size,
          output_dim=300,
          weights=[embedding_matrix],
          input_length=max_len,
          trainable=True)
```

- **Purpose:** Converts integer token sequences into dense vector representations using pretrained FastText embeddings (`wiki.simple.vec`, 300D).
- **Trainable:** Set to `True` to allow fine-tuning during training.
- **Weights:** Initialized using an embedding matrix built from FastText vectors.
- **Input Length:** Fixed using `max_len` (set during preprocessing and padding).

## 2. Bidirectional LSTM Layer

```
Bidirectional(LSTM(128, return_sequences=False))
```

- **Purpose:** Extracts contextual information from both the past and future (forward and backward) using 128 LSTM units.
- **Bidirectionality:** Helps capture sentiment and meaning more robustly than a unidirectional LSTM.
- **return\_sequences=False:** Only the final hidden state is used for classification (suitable for sentence-level prediction).

## 3. Dropout Layer

```
Dropout(0.5)
```

- **Purpose:** Prevents overfitting by randomly disabling 50% of the neurons during each training pass.

## 4. Dense Layer

```
Dense(32, activation='relu')
```

- **Purpose:** Acts as a non-linear transformation layer that learns intermediate features.
- **ReLU Activation:** Adds non-linearity and helps in learning complex relationships in data.

## 5. Output Layer

```
Dense(5, activation='softmax')
```

- **Purpose:** Predicts the probability distribution over 5 review classes (0 through 4), which are mapped to ratings 1 through 5 after prediction.
- **Softmax:** Ensures outputs represent class probabilities that sum to 1.

## Model Compilation

```
model.compile(loss='sparse_categorical_crossentropy',
              optimizer='adam',
              metrics=['accuracy'])
```

- **Loss Function:** `sparse_categorical_crossentropy` — used for multi-class classification with integer labels (instead of one-hot encoding).
- **Optimizer:** Adam — adaptive learning rate optimization suitable for NLP tasks.

- **Metric:** `accuracy` standard classification accuracy.

## Regularization: Early Stopping

```
EarlyStopping(monitor='val_loss', patience=2, restore_best_weights=True)
```

- **Why:** To prevent overfitting and long training time.
- **Mechanism:** Monitors the validation loss and stops training if it doesn't improve after 2 consecutive epochs.
- **Restores:** The model weights from the best epoch (lowest validation loss).

## Summary Table

| Layer Type     | Details   |
|----------------|---|
| Embedding      | FastText 300D, trainable, input length = <code>max_len</code> |
| BiLSTM         | 128 units, bidirectional                                      |
| Dropout        | 50%   |
| Dense (hidden) | 32 units, ReLU activation                                     |
| Output Dense   | 5 units, Softmax activation                                   |
| Loss Function  | Sparse Categorical Crossentropy                               |
| Optimizer      | Adam  |
| Early Stopping | Patience = 2, monitors validation loss                        |

## Word Embeddings: FastText (wiki.simple.vec)

### What Are Word Embeddings?

Word embeddings are dense vector representations of words where similar words have similar representations in vector space. Instead of treating words as discrete symbols, embeddings allow the model to **learn semantic relationships** between words.

### Which Embedding Is Used?

- **Type:** Pretrained **FastText** vectors
- **File Used:** `wiki.simple.vec`
- **Vector Dimension:** 300
- **Source:** Trained on **Simple English Wikipedia**
- **Format:** `.vec` file with lines like:

```
good 0.1234 -0.5432 ... 300 values
```

## Why FastText?

FastText is chosen over other embeddings like GloVe or Word2Vec for the following reasons:

| Feature                    | FastText vs Others  |
|----------------------------|---|
| <b>Subword Information</b> | FastText captures word parts (prefixes/suffixes), which improves handling of <b>rare</b> , <b>misspelled</b> , or <b>compound words</b> . Neither GloVe nor Word2Vec do this. |
| <b>Generalization</b>      | FastText can generate vectors for <b>OOV (Out-of-Vocabulary)</b> words by combining character n-grams — others return <code>null</code> .                                     |
| <b>Domain Relevance</b>    | The use of <b>Simple English Wikipedia</b> as the training corpus aligns well with product review language: clear, non-technical, and consumer-oriented.                      |
| <b>Fast to Load</b>        | <code>.vec</code> format is human-readable and easy to parse.   |
| <b>Compatibility</b>       | Easily converted into an embedding matrix compatible with TensorFlow/Keras models.  |

## Why LSTM? Why Bidirectional LSTM?

### What is an LSTM?

LSTM (**Long Short-Term Memory**) is a type of Recurrent Neural Network (RNN) that is designed to **learn patterns in sequential data**, such as text, time series, or speech.

It's particularly good at **remembering long-term dependencies** and **avoiding vanishing gradients**, which is a common issue with standard RNNs.

### How Does LSTM Work?

LSTM processes sequences **word by word** (or time step by time step) using a **memory cell** and three key gates:

| Gate               | Function  |
|--------------------|---|
| <b>Forget Gate</b> | Decides what information to discard from the cell state.                          |
| <b>Input Gate</b>  | Decides which new information to add to the cell state.                           |
| <b>Output Gate</b> | Decides what to output (i.e., what part of memory to pass to the next time step). |

### Internal Workflow at Each Step:

1. Takes a word/token as input.
2. Updates its **cell state** (memory) based on current input and past memory.
3. Uses the gates to control **information flow**.
4. Moves to the next token and repeats.

This ability to “remember” long-term dependencies helps LSTM capture **context** like:

- “I loved this product although delivery was late.”  
(final label = 4 or 5, despite late delivery)

## Why LSTM for This Task?

| Reason                                 | Explanation  |
|--|--|
| <b>Text is Sequential</b>              | Reviews are sentences or paragraphs — order of words matters. LSTMs are well-suited for this.      |
| <b>Contextual Understanding Needed</b> | To correctly predict a rating, the model needs to understand context (e.g., "not good" vs "good"). |
| <b>Better than CNNs for Long Texts</b> | LSTM captures longer dependencies; CNNs are better for fixed n-gram patterns.                      |
| <b>Robust to Varying Lengths</b>       | Can handle both short and long reviews effectively.  |

## Why *Bidirectional* LSTM?

A **Bidirectional LSTM** processes input in both directions:

- One LSTM processes text **left to right** (past → future)
- Another processes text **right to left** (future → past)

Then it **concatenates** both outputs to form a **full-context representation**.

### Benefits:

- Captures **dependencies from both ends** of the sentence.
- Improves performance in tasks like sentiment analysis, where key clues can appear at any point in the text.
- Handles patterns like:
  - "Not only was it slow, but it also broke" (need forward and backward context)

## Summary

| Component            | Reason for Selection   |
|----------------------|--|
| <b>FastText</b>      | Handles OOV words, uses subword units, Simple English Wiki trained |
| <b>LSTM</b>          | Ideal for sequential data, avoids vanishing gradients              |
| <b>Bidirectional</b> | Learns full context, improves accuracy in language tasks           |

## Strengths of LSTM

| <b>Strength</b>                        | <b>Why it Matters</b>  |
|--|--|
| <b>Captures long-term dependencies</b> | Understands meaning from distant words                       |
| <b>Avoids vanishing gradient</b>       | Learns more effectively over long text                       |
| <b>Works well with noisy data</b>      | Handles variations in text (spelling, style)                 |
| <b>Effective in NLP</b>                | State-of-the-art performance in sentiment, translation, etc. |
| <b>Supports variable input length</b>  | Works with both short and long reviews                       |

### When to Use LSTM

| <b>Use Case</b>                     | <b>Example</b>                             |
|-------------------------------------|--|
| <b>Sentiment analysis</b>           | Predicting positive/negative reviews       |
| <b>Text classification</b>          | Assigning topics or ratings to text        |
| <b>Sequence prediction</b>          | Predicting next word or sequence pattern   |
| <b>Speech and language modeling</b> | Translating sentences, voice recognition   |
| <b>Time series forecasting</b>      | Stock market trends, temperature over time |

### When *Not* to Use LSTM

| <b>Scenario</b>  | <b>Better Alternative</b>  |
|--|--|
| <b>Very large datasets with long training time</b>       | Transformer-based models (e.g., BERT, RoBERTa) are faster and scale better |
| <b>Highly structured, tabular data</b>                   | Use tree-based models like XGBoost or Random Forest                        |
| <b>Real-time deployment with low latency constraints</b> | LSTMs are slower than CNNs or smaller models                               |
| <b>Non-sequential tasks</b> (e.g., image classification) | Use CNNs instead   |

### Limitations of LSTM

| <b>Limitation</b>                              | <b>Impact</b>   |
|--|---|
| <b>Training can be slow</b>                    | Especially with long sequences                        |
| <b>High memory usage</b>                       | Due to recurrent operations and multiple gates        |
| <b>Hard to parallelize</b>                     | Each step depends on the previous one                 |
| <b>Struggles with extremely long sequences</b> | Even LSTM has limits; Transformers may perform better |

### Training Progress (Epoch-wise):

| Epoch | Accuracy      | Loss          | Val Accuracy  | Val Loss                                |
|-------|---------------|---------------|---------------|---|
| 1     | 43.26%        | 1.2804        | 53.43%        | 1.0782                                  |
| 2     | 56.02%        | 1.0246        | 54.51%        | 1.0714                                  |
| 3     | 59.41%        | 0.9499        | 54.67%        | 1.0857                                  |
| 4     | <b>63.17%</b> | <b>0.8698</b> | <b>53.93%</b> | <b>1.0938</b> ( <i>Early Stopping</i> ) |

### Interpretation

- **Model A** shows **progressive improvement** in training accuracy up to Epoch 4.
- Early stopping was triggered due to **no improvement in validation loss** beyond this point.
- **Overfitting is minimal**, as validation accuracy is consistently in the ~54% range.
- Despite being trained on a **perfectly balanced dataset**, the model generalizes **moderately well** even on validation samples