Logistic Regression for Imbalanced Dataset

Why Logistic Regression?

Logistic Regression is a well-established linear model widely used in text classification tasks due to its simplicity, efficiency, and solid theoretical grounding. It is especially effective when paired with TF-IDF vectorized data, as it can handle high-dimensional, sparse feature spaces often found in natural language processing (NLP) tasks. It also offers probabilistic outputs and interpretable coefficients.

We selected Logistic Regression as a baseline algorithm to classify product review ratings (1 to 5 stars) based on review text. It is scalable and offers good performance with appropriate preprocessing and hyperparameter tuning.

Model Training Logic

The pipeline followed was:

1. **Text Preprocessing**:
    - Lowercasing
    - Punctuation removal
    - Stopword removal (NLTK/spaCy)
    - Lemmatization
2. **Vectorization**: TF-IDF (`max_features=5000`)
3. **Model**: `LogisticRegression` trained using `GridSearchCV` for hyperparameter optimization.
4. **Data Split**: Stratified 80/20 train-test split to preserve rating distribution.
5. **Scoring Metric**: `f1_macro` — chosen to give equal weight to all classes, especially for imbalanced datasets.

Hyperparameter Tuning Details

```
param_grid = {
    'C': [0.01, 0.1, 1, 10],              # Regularization strength (inverse)
    'penalty': ['l2'],                    # Regularization type
    'solver': ['liblinear', 'lbfgs', 'saga'],  # Optimization algorithm
    'max_iter': [100, 200, 500]           # Max training iterations
}
```

- **C** controls regularization: lower = stronger regularization.

- **penalty** 'l2' helps prevent overfitting in high-dimensional text data.
- **solver** options tested for stability and speed.
- **GridSearchCV** used 3-fold cross-validation to select the best combination.

Evaluation Result

After hyperparameter tuning and training, the best model achieved:

Metric	Value
Accuracy	43%
Macro F1	42%
Support	2000 test samples

Detailed classification report:

	precision	recall	f1-score	support
1	0.58	0.35	0.44	200
2	0.40	0.25	0.31	300
3	0.36	0.42	0.39	500
4	0.43	0.54	0.48	600
5	0.51	0.45	0.47	400
accuracy			0.43	2000
macro avg	0.45	0.40	0.42	2000
weighted avg	0.44	0.43	0.42	2000

- **Class 4 & 5** performed better, likely due to more representation in training data.
- **Class 1 & 2** showed lower recall, suggesting the model struggles with minority classes in imbalanced settings

Interpretation

- The model demonstrates **moderate performance**, with higher confidence in predicting frequent ratings.
- It **fails to generalize** well on underrepresented ratings (like 1 and 2 stars), a common issue in imbalanced datasets.
- **Macro F1-score** of 0.42 highlights the need for more advanced balancing or feature techniques for improvement.

When to Use Logistic Regression

- High-dimensional, sparse feature space (like TF-IDF)
- Need for interpretable models
- Binary or multiclass classification with linearly separable data
- As a **baseline model** for text classification tasks

Limitations

- Assumes **linear decision boundaries**
- Sensitive to **correlated features**
- Struggles with **imbalanced classes** without tuning or sampling
- Less powerful than tree-based models or deep learning in non-linear settings

Confusion matrix